

Into. to Big-Data | OPPE (25-08-24) | report

By Chandresh J. Sutariya (21f3001415)

Report | Problem Statement:

Problem Statement:

The station master at every train station is tasked with continuously assessing whether the number of platforms available in that station at any point in time during the day is sufficient for the trains expected to stop at that station. Each station master is looking for help to solve for this challenge by having at his/her fingerstep the total number of trains at that station in any 20 minute window. It is assumed that a typical train needs about 20 minutes at a station at the maximum for loading & unloading before moving on or being moved to the shed.

Your task is to compute the 20-minute rolling count of trains per station so that every 5 minutes, every station master that has any train stopped at their station gets this count information for their purposes. You are also required to determine per station the maximum of these counts across the entire time period of the data set.

Report | Step

1. setup kafka:
 - Zookeeper, kafka server
 2. Create topic with name *Test-5*
 3. Start producer file *prdcr.py* on kafka VM
 4. Start consumer file *consmr.py* on DataProc Cluster VM
-

Batch: 1

StationCode	window	TrainsOnStation
JTJ	{2024-08-25 02:40:00, 2024-08-25 03:00:00}	2
GR	{2024-08-25 17:40:00, 2024-08-25 18:00:00}	1
BXL	{2024-08-25 22:20:00, 2024-08-25 22:40:00}	1
DVL	{2024-08-25 23:00:00, 2024-08-25 23:20:00}	2
VRL	{2024-08-25 17:00:00, 2024-08-25 17:20:00}	2
KNW	{2024-08-25 01:00:00, 2024-08-25 01:20:00}	3
AUBR	{2024-08-25 08:00:00, 2024-08-25 08:20:00}	1
SBM	{2024-08-25 05:20:00, 2024-08-25 05:40:00}	1
KUN	{2024-08-25 03:20:00, 2024-08-25 03:40:00}	1
JM	{2024-08-25 10:20:00, 2024-08-25 10:40:00}	1
FGR	{2024-08-25 17:00:00, 2024-08-25 17:20:00}	1
FL	{2024-08-25 17:20:00, 2024-08-25 17:40:00}	1
KVZ	{2024-08-25 18:40:00, 2024-08-25 19:00:00}	1
BSL	{2024-08-25 21:20:00, 2024-08-25 21:40:00}	1
KBY	{2024-08-25 07:00:00, 2024-08-25 07:20:00}	1
CAN	{2024-08-25 23:40:00, 2024-08-26 00:00:00}	1
TDD	{2024-08-25 20:00:00, 2024-08-25 20:20:00}	1
HJLI	{2024-08-25 05:40:00, 2024-08-25 06:00:00}	1
MZP	{2024-08-25 23:00:00, 2024-08-25 23:20:00}	1
MULK	{2024-08-25 06:00:00, 2024-08-25 06:20:00}	1

only showing top 20 rows

[Consumer clientId=consumer-spark-kafka-source-65227384-1ea9-4c5b-9cd8-9f2ac77a159b-365556117-driver-0-1, groupId=spark-kafka-source-65227384-1ea9-4c5b-9cd8-9f2ac77a159b-365556117-driver-0] Resetting offset for partition test-4-0 to offset 182147.

For the last batch, the offset will be set to total number of records that were there in the kafka-topic.

So, here it-is 1,82,147. And while I was preprocessing the data, the data after the preprocessing were also 182147. So that's one check!

df.count()

182147

```

1 from pyspark.sql import SparkSession
2 import os
3 import sys
4
5 from pyspark.ml import Pipeline
6 from pyspark.ml.feature import SQLTransformer
7
8
9 os.environ['PYSPARK_DRIVER_PYTHON'] = sys.executable
10 os.environ['PYSPARK_PYTHON'] = sys.executable
11
12 #####
13 # global var.
14 want_to_publish_data_on_kafak = True
15 file_path = '/home/chandreshjsutariya/Train_details_22122017.csv'
16 rows_to_send_count = 0
17 rows_to_send = 2
18 count=0
19 # Define Kafka producer parameters
20 kafka_bootstrap_servers = "35.231.64.61:9092"
21 kafka_topic = "test-2"
22 #####
23
24 spark = (
25     SparkSession.builder
26         .master("local")
27         .appName("ass-8")
28         .getOrCreate()
29 )
30 spark.conf.set("spark.sql.execution.arrow.pyspark.enabled", "true")
31
32 trains = spark.read.format('csv').load(file_path, header=True)
33 trains.show()

```

```

34
35
36 remove_null = SQLTransformer(statement="SELECT * FROM __THIS__ \
37 | | | | | | | | | | | | | | | | WHERE 'Arrival Time' IS NOT NULL AND 'Departure Time' IS NOT NULL")
38 trains = remove_null.transform(trains)
39
40 trains.count()
41 #186124
42
43 from pyspark.sql.functions import col, expr
44 # from kafka import KafkaProducer
45 import json
46
47 trains.count()
48
49 trains = trains.filter((col("Arrival time") != "00:00:00") & \
50 | | | | | | | | | | | | | | (col("Departure Time") != "00:00:00"))
51 trains.count()
52
53 trains = trains.withColumn("AT", expr(\
54 |   "to_timestamp(`Arrival time`, 'HH:mm:ss')"\
55 | )
56 | )
57 trains = trains.withColumn("DT", expr(\
58 |   "to_timestamp(`Departure Time`, 'HH:mm:ss')"\
59 | )
60 | )
61 trains.count()
62

```

```

62
63 trains = trains.withColumn("TimeOnPltf",
64                             expr("CASE WHEN `DT` >= `AT` " +
65                                   "THEN (unix_timestamp(`DT`) - unix_timestamp(`AT`)) / 60 " +
66                                   "ELSE (86400 - unix_timestamp(`AT`)) / 60 + (unix_timestamp(`DT`) / 60) END")
67 )
68 # trains.show()
69
70 # drop rows with null values in platformstaymin col
71 trains = trains.dropna(subset=["TimeOnPltf"])
72 trains.count()
73
74 # function to send message to kafka using kafka producer
75 from kafka import KafkaProducer
76
77 def send_to_kafka_partition(iter):
78     producer = KafkaProducer(bootstrap_servers = kafka_bootstrap_servers,
79                               value_serializer = lambda x: json.dumps(x).encode('utf-8'))
80     for row in iter:
81         message = {
82             "StationCode": row['Station Code'],
83             "AT": row['Arrival time'],
84             "DT": row['Departure Time'],
85             "TimeOnPltf": row['TimeOnPltf']
86         }
87         producer.send(kafka_topic, value=message)
88
89     producer.flush()
90

```

```

from pyspark.sql import SparkSession
from pyspark.sql.functions import window, col

spark = SparkSession\
    .builder\
    .appName("oppe-24")\
    .getOrCreate()

#####
# global var.
want_to_start_consumer = True
kafka_bootstrap_servers = '35.231.64.61:9092'
kafka_topic = 'test-5'
#####

df = spark \
    .readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", kafka_bootstrap_servers) \
    .option("subscribe", kafka_topic) \
    .option("startingOffsets", "earliest")\
    .load()\
    .selectExpr("CAST(value AS STRING) as value")

json_schema = "StationCode STRING, AT STRING, DT STRING, TrainName STRING, TimeOnPltf DOUBLE"

parsed_df = df.selectExpr("from_json(value, '{}') as data".format(json_schema)) \
    .select("data.*")\
    .filter(col("TimeOnPltf").isNotNull())

parsed_df = parsed_df.withColumn("AT", col("AT").cast("timestamp"))

parsed_df = parsed_df.withColumn("AT", col("AT").cast("timestamp"))

windowed_df = parsed_df.groupBy("StationCode", window("AT", "20 minutes", "20 minutes"
))\
    .count()\
    .withColumnRenamed("count", "TrainsOnStation")

query = windowed_df.writeStream\
    .outputMode("complete")\
    .format("console")\
    .option("truncate", "false")\
    .trigger(processingTime="5 seconds")\
    .start()

query.awaitTermination()
spark.stop()

```